

```
In [1]: import random
import numpy as np
from data.process import get_CIFAR10_data, get_MUSHROOM_data
from scipy.spatial import distance
from models import Perceptron, SVM, Softmax, Logistic
from kaggle_submission import output_submission_csv
%matplotlib inline
```

Loading CIFAR-10

In the following cells we determine the number of images for each split and load the images.
TRAIN_IMAGES + VAL_IMAGES = (0, 50000) , TEST_IMAGES = 10000

```
In [2]: # You can change these numbers for experimentation
# For submission we will use the default values
TRAIN_IMAGES = 40000
VAL_IMAGES = 10000
```

```
In [3]: data = get_CIFAR10_data(TRAIN_IMAGES, VAL_IMAGES)
X_train_CIFAR, y_train_CIFAR = data['X_train'], data['y_train']
X_val_CIFAR, y_val_CIFAR = data['X_val'], data['y_val']
X_test_CIFAR, y_test_CIFAR = data['X_test'], data['y_test']
n_class_CIFAR = len(np.unique(y_test_CIFAR))
```

Convert the sets of images from dimensions of (N, 3, 32, 32) -> (N, 3072) where N is the number of images so that each 3x32x32 image is represented by a single vector.

```
In [4]: X_train_CIFAR = np.reshape(X_train_CIFAR, (X_train_CIFAR.shape[0], -1))
X_val_CIFAR = np.reshape(X_val_CIFAR, (X_val_CIFAR.shape[0], -1))
X_test_CIFAR = np.reshape(X_test_CIFAR, (X_test_CIFAR.shape[0], -1))
```

```
In [5]: print('Train data shape: ', X_train_CIFAR.shape)
print('Train labels shape: ', y_train_CIFAR.shape)
print('Validation data shape: ', X_val_CIFAR.shape)
print('Validation labels shape: ', y_val_CIFAR.shape)
print('Test data shape: ', X_test_CIFAR.shape)
print('Test labels shape: ', y_test_CIFAR.shape)
print("Number of class: ", n_class_CIFAR)
```

```
Train data shape: (40000, 3072)
Train labels shape: (40000,)
Validation data shape: (10000, 3072)
Validation labels shape: (10000,)
Test data shape: (10000, 3072)
Test labels shape: (10000,)
Number of class: 10
```

Loading Mushroom

In the following cells we determine the splitting of the mushroom dataset.
TRAINING + VALIDATION = 0.8, TESTING = 0.2

```
In [6]: # TRAINING = 0.6 indicates 60% of the data is used as the training dataset.
VALIDATION = 0.2
```

```
In [7]: data = get_MUSHROOM_data(VALIDATION)
X_train_MR, y_train_MR = data['X_train'], data['y_train']
X_val_MR, y_val_MR = data['X_val'], data['y_val']
X_test_MR, y_test_MR = data['X_test'], data['y_test']
n_class_MR = len(np.unique(y_test_MR))

print("Number of train samples: ", X_train_MR.shape[1])
print("Number of val samples: ", X_val_MR.shape[1])
print("Number of test samples: ", X_test_MR.shape[1])
print("Number of class: ", n_class_MR)
```

```
Number of train samples: 22
Number of val samples: 22
Number of test samples: 22
Number of class: 2
```

Get Accuracy

This function computes how well your model performs using accuracy as a metric.

```
In [8]: def get_acc(pred, y_test):
return np.sum(y_test==pred)/len(y_test)*100
```

Perceptron

Perceptron has 2 hyperparameters that you can experiment with:

- Learning rate** - controls how much we change the current weights of the classifier during each update. We set it at a default value of 0.5, but you should experiment with different values. We recommend changing the learning rate by factors of 10 and observing how the performance of the classifier changes. You should also try adding a **decay** which slowly reduces the learning rate over each epoch.
- Number of Epochs** - An epoch is a complete iterative pass over all of the data in the dataset. During an epoch we predict a label using the classifier and then update the weights of the classifier according to the perceptron update rule for each sample in the training set. You should try different values for the number of training epochs and report your results.

You will implement the Perceptron classifier in the **models/Perceptron.py**

The following code:

- Creates an instance of the Perceptron classifier class
- The train function of the Perceptron class is trained on the training data
- We use the predict function to find the training accuracy as well as the testing accuracy

Train Perceptron on CIFAR

```
In [9]: lr = 0.1
n_epochs = 100

percept_CIFAR = Perceptron(n_class_CIFAR, lr, n_epochs)
percept_CIFAR.train(X_train_CIFAR, y_train_CIFAR)
```

```
In [10]: pred_percept = percept_CIFAR.predict(X_train_CIFAR)
print('The training accuracy is given by: %f' % (get_acc(pred_percept, y_train_CIFAR)))
```

The training accuracy is given by: 43.862500

Validate Perceptron on CIFAR

```
In [11]: pred_percept = percept_CIFAR.predict(X_val_CIFAR)
print('The validation accuracy is given by: %f' % (get_acc(pred_percept, y_val_CIFAR)))
```

The validation accuracy is given by: 36.680000

Test Perceptron on CIFAR

```
In [12]: pred_percept = percept_CIFAR.predict(X_test_CIFAR)
print('The testing accuracy is given by: %f' % (get_acc(pred_percept, y_test_CIFAR)))
```

The testing accuracy is given by: 36.860000

Perceptron_CIFAR Kaggle Submission

Once you are satisfied with your solution and test accuracy, output a file to submit your test set predictions to the Kaggle for Assignment 1 CIFAR. Use the following code to do so:

```
In [13]: output_submission_csv('kaggle/perceptron_submission_CIFAR.csv', percept_CIFAR.predict(X_test_CIFAR))
```

Train Perceptron on Mushroom

```
In [14]: lr = 0.1
n_epochs = 100

percept_MR = Perceptron(n_class_MR, lr, n_epochs)
percept_MR.train(X_train_MR, y_train_MR)
```

```
In [15]: pred_percept = percept_MR.predict(X_train_MR)
print('The training accuracy is given by: %f' % (get_acc(pred_percept, y_train_MR)))
```

The training accuracy is given by: 80.118999

Validate Perceptron on Mushroom

```
In [16]: pred_percept = percept_MR.predict(X_val_MR)
print('The validation accuracy is given by: %f' % (get_acc(pred_percept, y_val_MR)))
```

The validation accuracy is given by: 78.523077

Test Perceptron on Mushroom

```
In [17]: pred_percept = percept_MR.predict(X_test_MR)
print('The testing accuracy is given by: %f' % (get_acc(pred_percept, y_test_MR)))
```

The testing accuracy is given by: 79.569231

Support Vector Machines (with SGD)

Next, you will implement a "soft margin" SVM. In this formulation you will maximize the margin between positive and negative training examples and penalize margin violations using a hinge loss.

We will optimize the SVM loss using SGD. This means you must compute the loss function with respect to model weights. You will use this gradient to update the model weights.

SVM optimized with SGD has 3 hyperparameters that you can experiment with :

- Learning rate** - similar to as defined above in Perceptron, this parameter scales by how much the weights are changed according to the calculated gradient update.
- Epochs** - similar to as defined above in Perceptron.
- Regularization constant** - Hyperparameter to determine the strength of regularization. In this case it is a coefficient on the term which maximizes the margin. You could try different values. The default value is set to 0.05.

You will implement the SVM using SGD in the **models/SVM.py**

The following code:

- Creates an instance of the SVM classifier class
- The train function of the SVM class is trained on the training data
- We use the predict function to find the training accuracy as well as the testing accuracy

Train SVM on CIFAR

```
In [18]: n_class_CIFAR = 10
lr = 0.01
n_epochs = 300
reg_const = 0.05

svm_CIFAR = SVM(n_class_CIFAR, lr, n_epochs, reg_const)
svm_CIFAR.train(X_train_CIFAR, y_train_CIFAR)

start training
```

```
In [19]: pred_svm = svm_CIFAR.predict(X_train_CIFAR)
print('The training accuracy is given by: %f' % (get_acc(pred_svm, y_train_CIFAR)))
```

start predicting
The training accuracy is given by: 33.725000

Validate SVM on CIFAR

```
In [20]: pred_svm = svm_CIFAR.predict(X_val_CIFAR)
print('The validation accuracy is given by: %f' % (get_acc(pred_svm, y_val_CIFAR)))
```

start predicting
The validation accuracy is given by: 32.100000

Test SVM on CIFAR

```
In [21]: pred_svm = svm_CIFAR.predict(X_test_CIFAR)
print('The testing accuracy is given by: %f' % (get_acc(pred_svm, y_test_CIFAR)))
```

start predicting
The testing accuracy is given by: 32.680000

SVM_CIFAR Kaggle Submission

Once you are satisfied with your solution and test accuracy output a file to submit your test set predictions to the Kaggle for Assignment 1 CIFAR. Use the following code to do so:

```
In [22]: output_submission_csv('kaggle/svm_submission_CIFAR.csv', svm_CIFAR.predict(X_test_CIFAR))

start predicting
```

Train SVM on Mushroom

```
In [23]: n_class_MR = 2
lr = 0.01
n_epochs = 300
reg_const = 0.05

svm_MR = SVM(n_class_MR, lr, n_epochs, reg_const)
svm_MR.train(X_train_MR, y_train_MR)

start training
```

```
In [24]: pred_svm = svm_MR.predict(X_train_MR)
print('The training accuracy is given by: %f' % (get_acc(pred_svm, y_train_MR)))
```

start predicting
The training accuracy is given by: 89.967173

Validate SVM on Mushroom

```
In [25]: pred_svm = svm_MR.predict(X_val_MR)
print('The validation accuracy is given by: %f' % (get_acc(pred_svm, y_val_MR)))
```

start predicting
The validation accuracy is given by: 89.784615

Test SVM on Mushroom

```
In [26]: pred_svm = svm_MR.predict(X_test_MR)
print('The testing accuracy is given by: %f' % (get_acc(pred_svm, y_test_MR)))
```

start predicting
The testing accuracy is given by: 89.107692

Softmax Classifier (with SGD)

Next, you will train a Softmax classifier. This classifier consists of a linear function of the input data followed by a softmax function which outputs a vector of dimension C (number of classes) for each data point. Each entry of the softmax output vector corresponds to a confidence in one of the C classes, and like a probability distribution, the entries of the output vector sum to 1. We use a cross-entropy loss on this softmax output to train the model.

Check the following link as an additional resource on softmax classification: <http://cs231n.github.io/linear-classify/#softmax>

Once again we will train the classifier with SGD. This means you need to compute the gradients of the softmax cross-entropy loss function according to the weights and update the weights using this gradient. Check the following link to help with implementing the gradient updates: <https://deeppoints.io/softmax-crossentropy>

The softmax classifier has 3 hyperparameters that you can experiment with :

- Learning rate** - As above, this controls how much the model weights are updated with respect to their gradient.
- Number of Epochs** - As described for perceptron.
- Regularization constant** - Hyperparameter to determine the strength of regularization. In this case, we minimize the L2 norm of the model weights as regularization, so the regularization constant is a coefficient on the L2 norm in the combined cross-entropy and regularization objective.

You will implement a softmax classifier using SGD in the **models/Softmax.py**

The following code:

- Creates an instance of the Softmax classifier class
- The train function of the Softmax class is trained on the training data
- We use the predict function to find the training accuracy as well as the testing accuracy

Train Softmax on CIFAR

```
In [27]: lr = 0.5
n_epochs = 300
reg_const = 0.05

softmax_CIFAR = Softmax(n_class_CIFAR, lr, n_epochs, reg_const)
softmax_CIFAR.train(X_train_CIFAR, y_train_CIFAR)

start training
```

```
In [28]: pred_softmax = softmax_CIFAR.predict(X_train_CIFAR)
print('The training accuracy is given by: %f' % (get_acc(pred_softmax, y_train_CIFAR)))
```

start predicting
The training accuracy is given by: 24.347500

Validate Softmax on CIFAR

```
In [29]: pred_softmax = softmax_CIFAR.predict(X_val_CIFAR)
print('The validation accuracy is given by: %f' % (get_acc(pred_softmax, y_val_CIFAR)))
```

start predicting
The validation accuracy is given by: 23.470000

Testing Softmax on CIFAR

```
In [30]: pred_softmax = softmax_CIFAR.predict(X_test_CIFAR)
print('The testing accuracy is given by: %f' % (get_acc(pred_softmax, y_test_CIFAR)))
```

start predicting
The testing accuracy is given by: 23.700000

Softmax_CIFAR Kaggle Submission

Once you are satisfied with your solution and test accuracy output a file to submit your test set predictions to the Kaggle for Assignment 1 CIFAR. Use the following code to do so:

```
In [31]: output_submission_csv('kaggle/softmax_submission_CIFAR.csv', softmax_CIFAR.predict(X_test_CIFAR))

start predicting
```

Train Softmax on Mushroom

```
In [32]: lr = 0.5
n_epochs = 300
reg_const = 0.05

softmax_MR = Softmax(n_class_MR, lr, n_epochs, reg_const)
softmax_MR.train(X_train_MR, y_train_MR)

start training
```

```
In [33]: pred_softmax = softmax_MR.predict(X_train_MR)
print('The training accuracy is given by: %f' % (get_acc(pred_softmax, y_train_MR)))
```

start predicting
The training accuracy is given by: 62.084530

Validate Softmax on Mushroom

```
In [34]: pred_softmax = softmax_MR.predict(X_val_MR)
print('The validation accuracy is given by: %f' % (get_acc(pred_softmax, y_val_MR)))
```

start predicting
The validation accuracy is given by: 60.430769

Testing Softmax on Mushroom

```
In [35]: pred_softmax = softmax_MR.predict(X_test_MR)
print('The testing accuracy is given by: %f' % (get_acc(pred_softmax, y_test_MR)))
```

start predicting
The testing accuracy is given by: 58.400000

Logistic Classifier

The Logistic Classifier has 2 hyperparameters that you can experiment with:

- Learning rate** - similar to as defined above in Perceptron, this parameter scales by how much the weights are changed according to the calculated gradient update.
- Number of Epochs** - As described for perceptron.

You will implement the Logistic Classifier in the **models/Logistic.py**

The following code:

- Creates an instance of the Logistic classifier class
- The train function of the Logistic class is trained on the training data
- We use the predict function to find the training accuracy as well as the testing accuracy

Training Logistic Classifier

```
In [36]: learning_rate = 0.1
n_epochs = 10

lr = Logistic(learning_rate, n_epochs)
lr.train(X_train_MR, y_train_MR)
```

```
In [37]: pred_lr = lr.predict(X_train_MR)
print('The training accuracy is given by: %f' % (get_acc(pred_lr, y_train_MR)))
```

The training accuracy is given by: 89.495281

Validate Logistic Classifier

```
In [38]: pred_lr = lr.predict(X_val_MR)
print('The validation accuracy is given by: %f' % (get_acc(pred_lr, y_val_MR)))
```

The validation accuracy is given by: 88.615385

Test Logistic Classifier

```
In [39]: pred_lr = lr.predict(X_test_MR)
print('The testing accuracy is given by: %f' % (get_acc(pred_lr, y_test_MR)))
```

The testing accuracy is given by: 89.538462

```
In [ ]:
```

